

PW: Is there a way to know which loss function one should be using in certain case?

Loss function is always application/problem specific.
We will look at a few of them in the course specific to classification and regression

Classification /
Regression


EL: when will this be?

The lecture video following the backprop is on loss functions. Then there will be demos, MAs and GAs where you will see them in code

Specific loss function for Classification will be revisited in Convolutional NN and for regression in Recurrent NN modules as well

EL: On convolutional models, will we be going through methods like image super sampling and neural texture compression? ^{resolution}

We will look at classification and segmentation model architectures. Super resolution is similar to segmentation with difference in how the reconstruction is used within the architecture.



There are some architectures that form fundamental building blocks (hence part of the intro course). Once you understand them you can carry forward this knowledge to understand other models. In that sense texture compression is too niche to focus on

